

Learned Gender Bias in Wikipedia Persists with Time, as measured by WEAT Scores

Blake Vente

Hunter College

Ralph.Vente09@myhunter.cuny.edu

Abstract

Word embeddings exhibit desirable properties when converting natural language to numerical vector representations. However, embeddings often internalize associations that parrot stereotypes pertaining to race, gender, and culture. Researchers have attempted to mitigate word embedding bias by altering the model after training, or changing the loss function, but to address bias comprehensively, Brunet et al. (2018) turn to the data where these biases originate. This work examines the differences in word embedding representations learned of Wikipedia a decade apart, finding minimal differences between them. This work finds that larger window sizes correlate with higher WEAT scores and higher performance on analogy and similarity datasets.

1 Introduction

Distributional semantic models represent words with fixed-dimensional vectors based on the how words are used in context of other words in large quantities of unstructured text (Lison and Kutuzov, 2017). Word embedding models in particular, represent words as low-dimensional vectors intended to capture functional and topical relations between words (Lison and Kutuzov, 2017).

1.1 Density of Word Representations

Classical approaches would represent words using one-hot encoding, with higher-dimensional vectors and binary elements. The one-hot representation of puppy would be a 1 in the ‘puppy’ dimension and zeros in all other components. Since vocabulary size is large even for moderately-sized corpora and since the dimensions of the vector depends on the vocabulary size, one-hot representations constrain the complexity of machine learning models

My source code associated with this paper is located at [rvente/word-embedding-bias](https://github.com/rvente/word-embedding-bias) on GitHub.




Figure 1: 50-dimensional vector representation of king from Jay Alammr’s *The Illustrated Word2vec*. The embedding is a dense vector of elements in the reals. The deepest red denotes components with higher magnitude, those close to 1.6 and the coolest blue denotes values close to -1.6 . One can see the similar components among all vectors (perhaps those encode a trait common to all entities, “humanity” in this case). One can observe the common components of ‘man’ and ‘woman’

where training complexity depends on the number of dimensions in the input. By contrast, word embeddings are fixed-dimensional representations with fixed dimension size of around 300 compared to the vocabularies of at least tens of thousands of words.

Word embeddings have a wide variety of applications for natural language processing tasks including part-of-speech tagging, syntactic parsing, and named entity recognition (Lison and Kutuzov, 2017). In general, word embeddings are broadly applicable for any machine learning system that operates on vectors, including deep learning models and . Independently of their training mechanisms, word embeddings reproduce the bias intrinsic to the data they were trained on.

1.2 Learning Semantic Representations

One key property of word embedding models is that that they can be trained on large corpora of unstructured text. This gives the key advantage that we can

retrain words representations as language continues to evolve. Two prominent families of models that exhibit semantic representations of words are Continuous Bag-of-Words (CBOW) models and Skip-gram models.

In CBOW models, context is used to predict a target word (Mikolov et al., 2013a). Formally, training example (x, y) is constructed by aggregating the representations of context words to form x , while the y . In particular, this “context” is defined as the k -word window about the target word w_i in document D . If $k = 2$, and the aggregation function is a simple sum operation, then $x = \Sigma(w_{i-2}, w_{i-1}, w_{i+1}, w_{i+2})$ and $y = w_i$. Then, i is increased, advancing the sliding window throughout the entire document, yielding further training examples. Emanating from this general method, further variations exist to enhance performance including weighting words by distance to the target.

In Skip-Gram models, the training principle is flipped: a target word is used to predict context words. The set of training examples would be $\{(w_i, w_{i-2}), (w_i, w_{i-1}), (w_i, w_{i+1}), (w_i, w_{i+2})\}$ (Mikolov et al. (2013a)). Then, in general, a dense neural network is chosen to learn on these examples and a loss function encapsulates the desire to “maximize similarity between semantically similar words”. A list of practical differences between these two training schemes is presented in Mikolov et al. (2013a), but is elided here.

1.3 Word Vector Arithmetic

Let $T : V \rightarrow E$ denote the embedding operation of word v from one-hot encoded (large) vocabulary V into embedding space E with finite dimension $\dim E$. Then we can define arithmetic on these vector representations. Doing so allows models to perform analogies.

$$T(\text{king}) - T(\text{man}) + T(\text{woman}) \approx T(\text{queen})$$

And as a form of transfer learning, it acts as a way of imbuing models with knowledge about syntactic and semantic relationships between words.

2 Word Embedding Bias

Reiterating $T : V \rightarrow E$ denotes the embedding operation of word v from vocabulary V into embedding space E with finite dimension $\dim E$, then we produce the well-documented and infamous expression $T(\text{computer programmer}) - T(\text{man}) + T(\text{woman}) \approx T(\text{homemaker})$ from (Bolukbasi

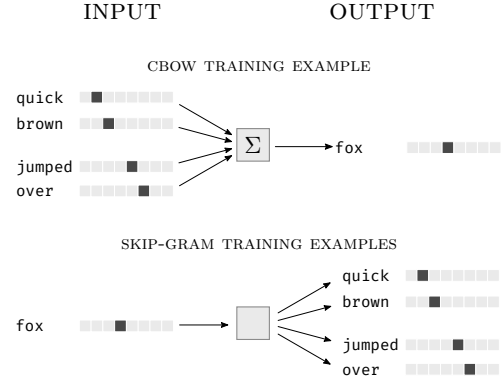


Figure 2: This figure, inspired by one in Mikolov et al. (2013a) illustrates the one-hot representations of words from the sentence “The quick brown fox jumped over ...” being summed to form a CBOW training example, while multiple examples are formed with the skipgram model. The squares denote components in the one-hot vector, with black denoting the 1 and light grey denoting the 0 elements. $\text{fox} = w_4$ in this imaginary corpus, making $\text{brown} = w_3$ and so on.

et al., 2016), which clearly exhibits the gender bias internalized by the learning algorithm that produced it. Hereafter, I denote the vector representation of the word in upright boldface.

The presence of these dubious associations is not unique to one architecture or one configuration of hyper-parameters (Brunet et al., 2018). Thus, the inclusion of a word embedding step may compromise the integrity of downstream machine learning operations by injecting or accentuating this bias, making the process unsuitable for a wide range of applications. In sensitive domains such as granting loans, such model behavior may be illegal.

3 Prior Work

Many researchers have attempted to reduce bias in these embedding models, with limited success. Caliskan et al. (2017). These works can be partitioned into two categories. First, there are those that alter the training process in some capacity. Then, there are those “post-processing” methods that alter vector representations at the end of training. Two such works follow: they define bias and attempt to minimize it without compromising performance on analogy tasks.

For example Bolukbasi et al. (2016) formulate that the gender bias in a non-gendered word can be quantified its scalar projection on the $\vec{\text{he}} - \vec{\text{she}}$ axis.

They zero out the first principal component in the gender direction. Gonen and Goldberg (2019) note that although the work was "extensive, thoughtful, and rigorous", the approach of Bolukbasi et al. (2016) is inherently limited as the chosen definition of bias is hand-selected.

By contrast, Zhao et al. (2018) also attempt to mitigate historical biases in word embeddings, but they do so by altering the loss function of GloVe to concentrate onto the last element the component of the embedding most correlated with gender. Then, at inference time the last element of the embedding is truncated away and thus discarded, "encouraging" the representations of non-gendered words to be orthogonal to the gender direction. Gonen and Goldberg (2019) opine that this method carries the right intuition – the alterations the the model needs to happen at training time, but that the execution has limited efficacy. In fact, "indirect bias" is still very obvious even when "direct bias" is mitigated. They demonstrate that even simple models can still learn the latent biases in the word embedding. This suggests that the components of the vector that encode stereotypical notions of gender are distributed as a linear combination of many components, even if they are orthogonal to a primary "gender dimension."

The general critique of both methods is that the definition of bias relating to distance to the gender direction Gonen and Goldberg (2019) believe that the correct method for removing gender bias, at a minimum, alters the training process.

Furthermore, Gonen and Goldberg (2019) remark that debiasing methods that don't take the data into consideration merely "cover-up" the biases without addressing the full associations themselves.

Historical Biases in Word Embedding models are pervasive: The stereotypes reproduced by word embeddings are not just limited to gender, but also extend to race and culture (Caliskan et al., 2017).

4 WEAT: Operationalizing Bias

As the desire to de-bias models gained traction, there was little consensus on how to quantify bias, making it difficult to compare debiasing methods. Thus, Islam et al. (2016) developed the Word Embedding Association Test (WEAT) to quantify how word embeddings capture empirical information about the world from text corpora (Islam et al., 2016, 8). WEAT score was modeled after the Im-

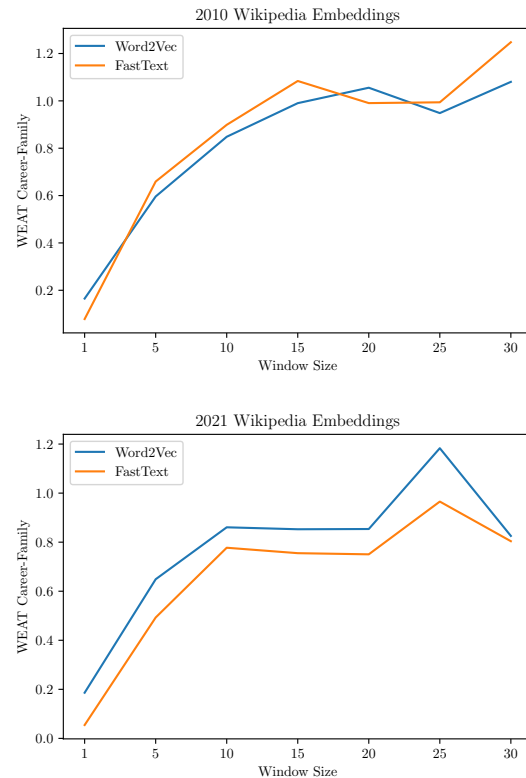


Figure 3: This comparison between 2010 and 2021 Wikipedia Embeddings shows that they behave roughly the same: up to a point, larger window sizes produce embedding models with higher WEAT scores. It does not suggest that one year's data is any more biased than another's: learned gender bias in Wikipedia as measured by WEAT persists with time, however window size increases bias as measured by WEAT

licit Association Test (IAT) as a source of documented human biases (Islam et al., 2016, 2). Intuitively, WEAT can be understood as a generalization of the use of word embedding models to perform analogies. They first define $s(w, A, B) =$

$$\frac{\text{mean}_{a \in A} \cos(\hat{w}, \hat{a}) - \text{mean}_{b \in B} \cos(\hat{w}, \hat{b})}{\text{std}_{x \in A \cup B} \cos(\hat{w}, \hat{x})}$$

and then $S(X, Y, A, B) = \sum_{x \in X} s(x, A, B) - \sum_{y \in Y} s(y, A, B)$. That is, S measures how close the associations between the target and attribute are. The more similar two sets of words are, the higher the WEAT score will be between them. The minimum WEAT score is -2 and the maximum is 2. The higher WEAT score is, the higher the agreement between the embedding model and we WEAT axis.¹ WEAT has been performed along several axes with words corresponding to the IAT as mentioned prior. I use the Career-Family and Math-Art axes, but model behavior was consistent between them. It is useful to note that while the remainder of this work uses WEAT as a stand-in for historical bias, it is only one approach of many that quantifies these stereotypes.

5 Experimental Setup

To examine the behavior of historical biases in word embeddings over time, I train GloVe and FastText word embeddings using the Gensim library by Řehůřek and Sojka (2010).

How does model behavior influenced by context window size? I declare window sizes of $W = \{1, 5, 10, 15, 20, 25, 30\}$ about the target word. Narrower window sizes usually encode syntactic relationships while longer sizes usually perform better on semantic relationships such as relational analogies Clark et al. (2013).

I declare my architectures $A = \{\text{FastText}, \text{Word2Vec}\}$; and corpora $C = \{\text{WikiSmall21}, \text{WikiSmall10}\}$ so named because they include only the first 330 MB of Wikipedia due to computational constraints. For preprocessing, I used Attardi (2014) to process the compressed archive of English Wikipedia. In all, I performed experiments $W \times A \times C$ to examine the relationships between model behaviors and window size, and how those behaviors change with time. Each year’s train-evaluate loop took about 90

¹To oversimplify, the higher the WEAT score is, the more it exhibits the particular bias along that axis.

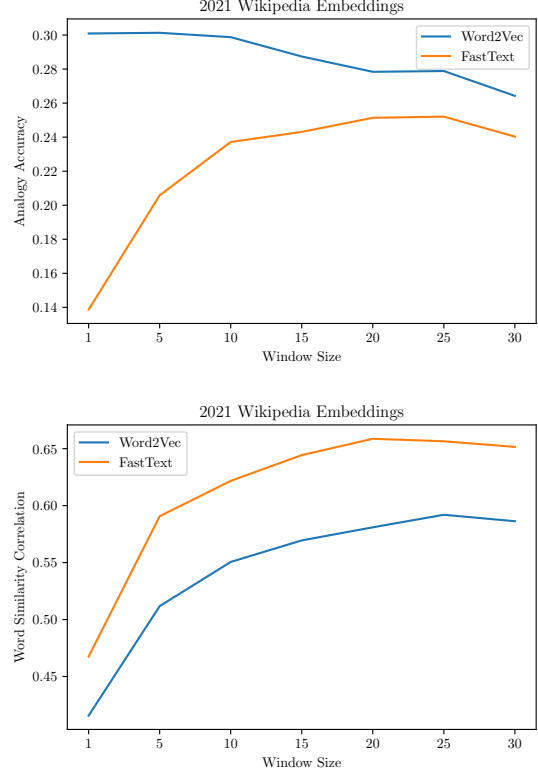


Figure 4: The Analogy task (above) peaks early for Word2Vec, yielding a negative correlation with window size as more “distracting” words are introduced into the context window. FastText does not exhibit this property increasing before a plateau. Word Similarity (below) also responds positively with window size, also with diminishing returns as window sizes get large.

minutes. Some portions are parallelized in Gensim, while other portions run on a single thread.

6 Evaluation

Perhaps the most remarkable part of word embeddings trained in these ways is that these vector representations can perform on tasks they were not trained directly on, including word similarity. For example, using Pearson or Spearman correlation, one can quantify how accurately the cosine of embedding representations predicts degree of similarity WordSim353 from (Agirre et al., 2009), with manually annotated word pairs quantifying human-percieved “similarity” between the words.

For evaluation on analogies, one can use (Mikolov et al., 2013b) among many choices. This has 19,544 instances and it is unbalanced, with 8,869 semantic and 10,675 syntactic questions, with between 20-70 pairs per category. Since country to capital relations comprise over 50 percent of

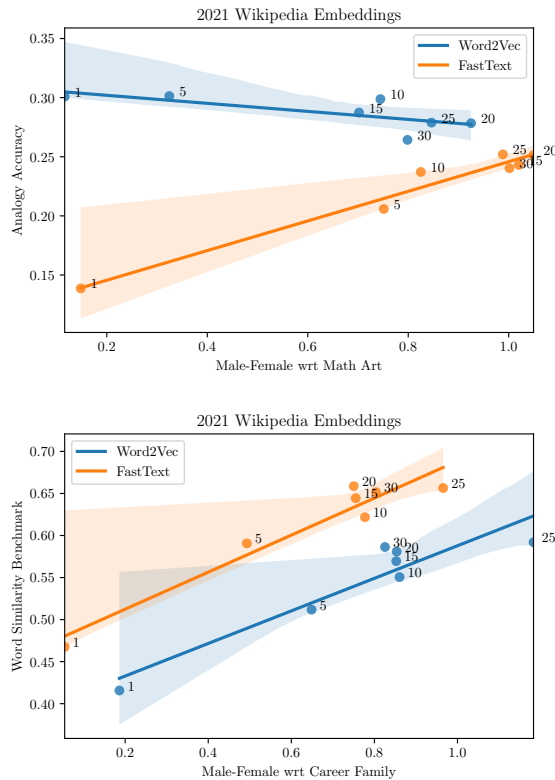


Figure 5: This is the scatter plot of Career-Family WEAT score against Analogy accuracy (above) and Word Similarity (below) this positive correlation suggests that models that have higher WEAT scores end to be the ones that perform better on convention benchmarks as well.

all semantic questions and in the 330 MB of training data, we can expect that there are many missing words that count against the accuracy. There is an opportunity for future work to compare different word sets.

7 Results

Figure 3 suggests that as window size increases, so does WEAT score. It suggests no substantive differences due to the passage of time in this sample of the Wikipedia corpus. Figure 4 shows the result of window size on conventional performance metrics. Figure 5 relates the performance on conventional metrics to the Career-Family WEAT metric.

8 Implications

This reinforces the notion that word associations that reproduce historical biases and stereotypes are learned the same way as factual word associations. However, it is possible that there are analogies in the dataset that may themselves exhibit stereotypes.

In this case, a “lesser-biased” vector representation may still exhibit a positive correlation as observed here. This suggests the need to verify that the analogy set itself before proceeding further with this claim.

9 Limitations and Future Work

1. Chief on the list of limitations is the relatively small data size used and that training only occurred for 1 epoch. Rather than the first 330 MB of data, future researchers can extend this to larger values keeping in mind the added training time.
2. WEAT has come under question as a valid metric for quantifying bias. Future work might evaluate WEAT in conjunction with a variety of other methods of quantifying historical bias.
3. There are many current attitudes discussing which domains are appropriate for debiasing, and whether some domains would benefit from “faithful” representations of the world. Future work might choose to engage with this.
4. There is work modeling analogy as probabilistic grammar and using statistical methods to learn analogies. Future work might compare these to word embeddings for these particular tasks.
5. Sampling window sizes between 1, 5, and 10 might reveal trends that are not clearly visible with the current set of window sizes.

10 Acknowledgements

Thank you Dr Anita Raja for helping me shape this research question and for extensive feedback at every step of this process. Furthermore, thank you Dr Sarah Ita Levitan for providing insight into practical considerations when incorporating a word embedding step and associated discussion and sparking my original interest in the topic.

References

- Eneko Agirre, Enrique Alfonseca, Keith Hall, Jana Kravalova, Marius Pasca, and Aitor Soroa. 2009. A study on similarity and relatedness using distributional and wordnet-based approaches.
- Giuseppe Attardi. 2014. Wikiextractor. <https://github.com/attardi/wikiextractor>.

- Tolga Bolukbasi, Kai-Wei Chang, James Y. Zou, Venkatesh Saligrama, and Adam Kalai. 2016. [Man is to computer programmer as woman is to home-maker? debiasing word embeddings.](#) *CoRR*, abs/1607.06520.
- Marc-Étienne Brunet, Colleen Alkalay-Houlihan, Ashton Anderson, and Richard S. Zemel. 2018. [Understanding the origins of bias in word embeddings.](#) *CoRR*, abs/1810.03611.
- Aylin Caliskan, Joanna J Bryson, and Arvind Narayanan. 2017. Semantics derived automatically from language corpora contain human-like biases. *Science*, 356(6334):183–186.
- Alexander Clark, Chris Fox, and Shalom Lappin. 2013. *The handbook of computational linguistics and natural language processing*. John Wiley & Sons.
- Hila Gonen and Yoav Goldberg. 2019. [Lipstick on a pig: Debiasing methods cover up systematic gender biases in word embeddings but do not remove them.](#) *CoRR*, abs/1903.03862.
- Aylin Caliskan Islam, Joanna J. Bryson, and Arvind Narayanan. 2016. [Semantics derived automatically from language corpora necessarily contain human biases.](#) *CoRR*, abs/1608.07187.
- Pierre Lison and Andrey Kutuzov. 2017. Redefining context windows for word embedding models: An experimental study. *arXiv preprint arXiv:1704.05781*.
- Tomas Mikolov, Quoc V. Le, and Ilya Sutskever. 2013a. [Exploiting similarities among languages for machine translation.](#)
- Tomas Mikolov, Ilya Sutskever, Kai Chen, Greg Corrado, and Jeffrey Dean. 2013b. [Distributed representations of words and phrases and their compositionality.](#)
- Radim Řehůřek and Petr Sojka. 2010. Software Framework for Topic Modelling with Large Corpora. In *Proceedings of the LREC 2010 Workshop on New Challenges for NLP Frameworks*, pages 45–50, Valletta, Malta. ELRA.
- Jieyu Zhao, Yichao Zhou, Zeyu Li, Wei Wang, and Kai-Wei Chang. 2018. [Learning gender-neutral word embeddings.](#)